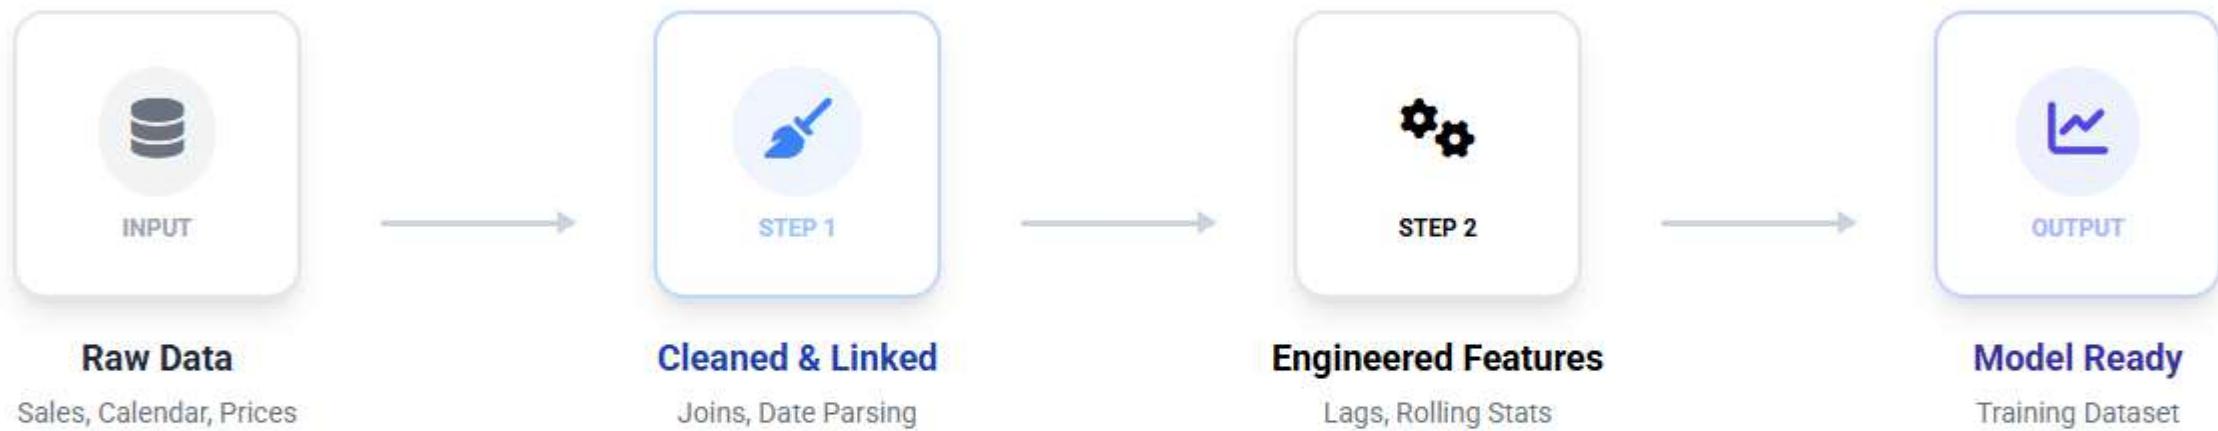


# 1. Role of Data Preprocessing

## Overview of the Forecasting Pipeline

Pipeline Overview



### What the pipeline does

Systematically **cleans** errors, **links** separate data sources, and **enriches** the raw logs with mathematical patterns needed for forecasting.

### Why it is required

Raw retail data is noisy and disconnected. Preprocessing **reduces noise**, aligns timelines, and adds context to **prevent data leakage**.

## 2. Raw Datasets Used

Overview of the primary data sources required for retail forecasting



### Sales Data

#### KEY COLUMNS

- ▢ item\_id
- ▢ store\_id
- ▢ day (d\_1, d\_2...)
- ▢ units\_sold

#### PURPOSE

*Defines historical demand: what was sold, where, and how much.*



### Calendar Data

#### KEY COLUMNS

- # day (d\_1...)
- date (YYYY-MM-DD)
- ▢ weekday
- ★ event\_name / type

#### PURPOSE

*Provides temporal context: connects abstract "Day 1" to real dates and holidays.*



### Sell Prices

#### KEY COLUMNS

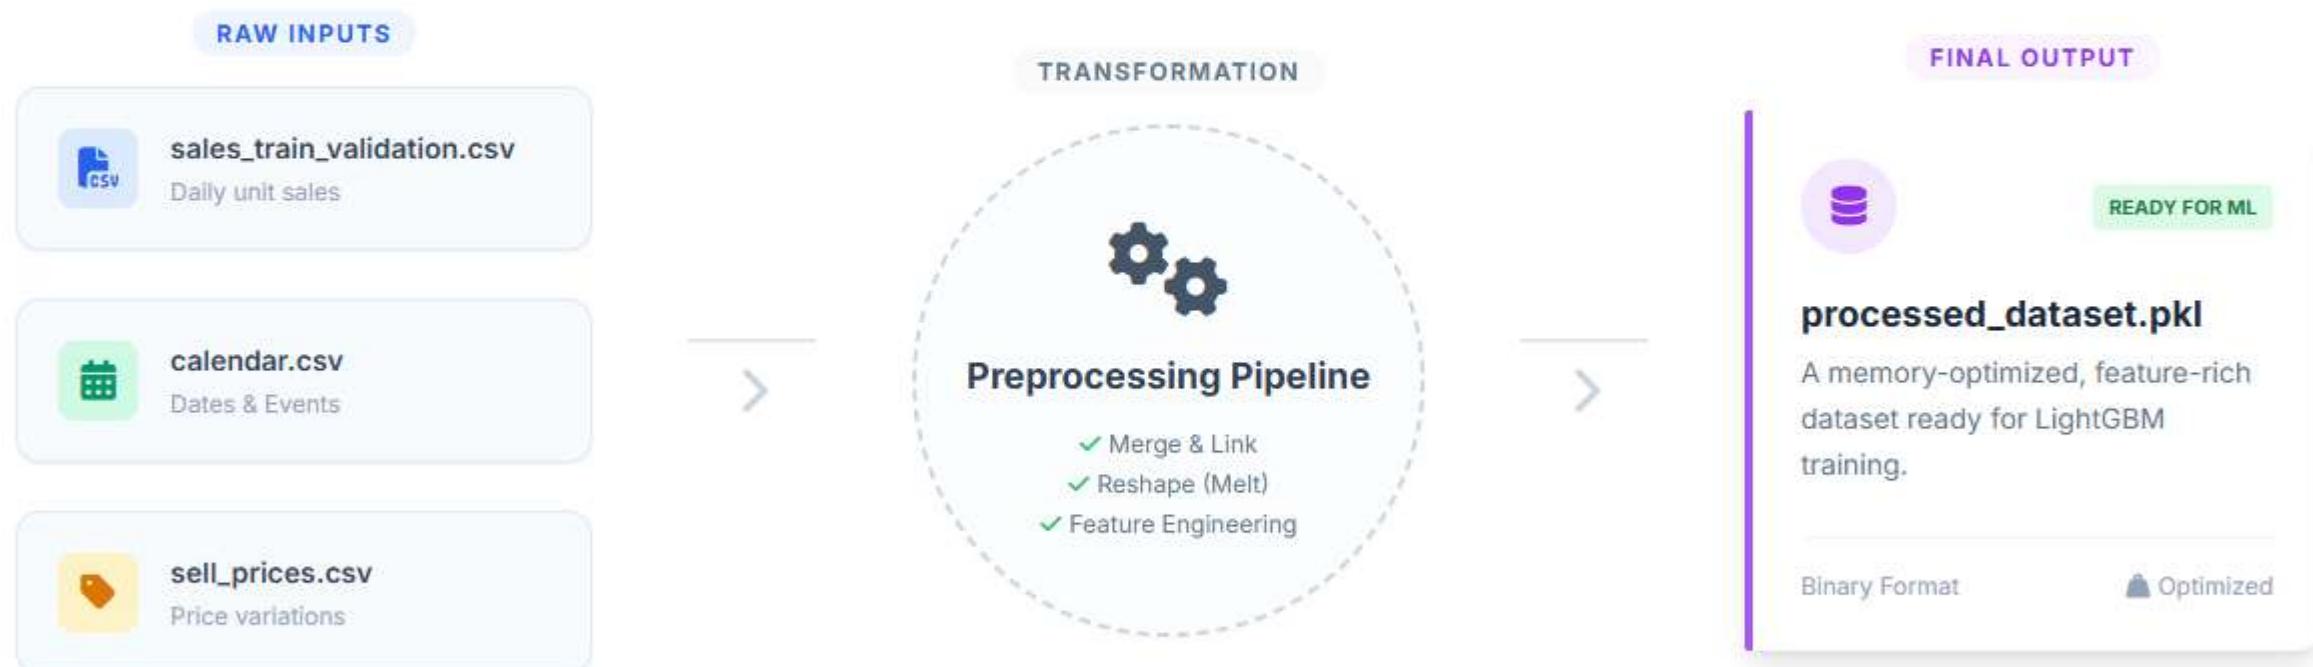
- ▢ item\_id
- ▢ store\_id
- wm\_yr\_wk (week id)
- § sell\_price

#### PURPOSE

*Tracks value: Price changes often drive demand shifts and promotions.*

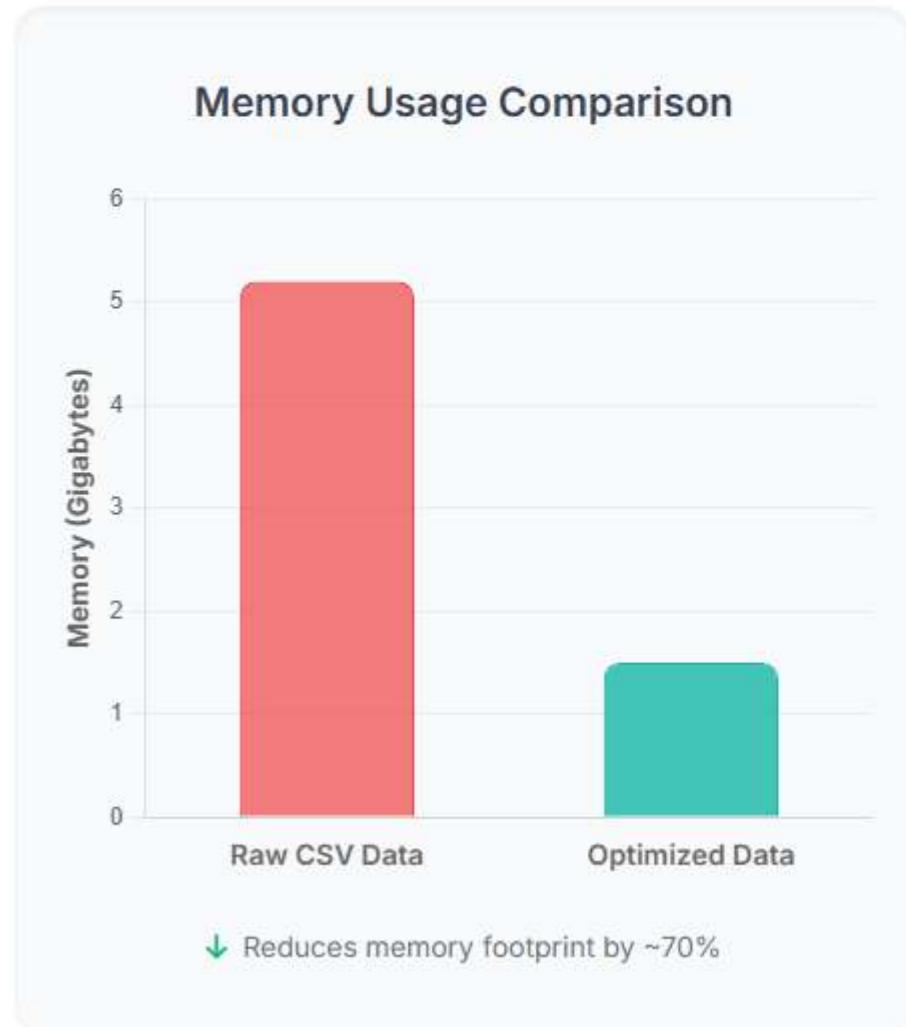
# Data Flow Overview

End-to-end transformation pipeline from raw sources to model-ready features



### 3. Why Large Retail Data Needs Memory Optimization

Handling millions of rows efficiently before processing



#### The Challenge

Retail datasets contain millions of transactions (rows) multiplied by thousands of items. Loading raw CSV files directly often causes **Out of Memory (OOM)** errors on standard computers.

#### Optimization Actions



##### Downcast Numeric Types

Convert 64-bit numbers to 16/32-bit (e.g., `float64 → float32`). We don't need 15 decimal places for sales counts.



##### Categorical Encoding

Store repeated strings (like "State\_CA") as compact integers internally, saving massive string overhead.

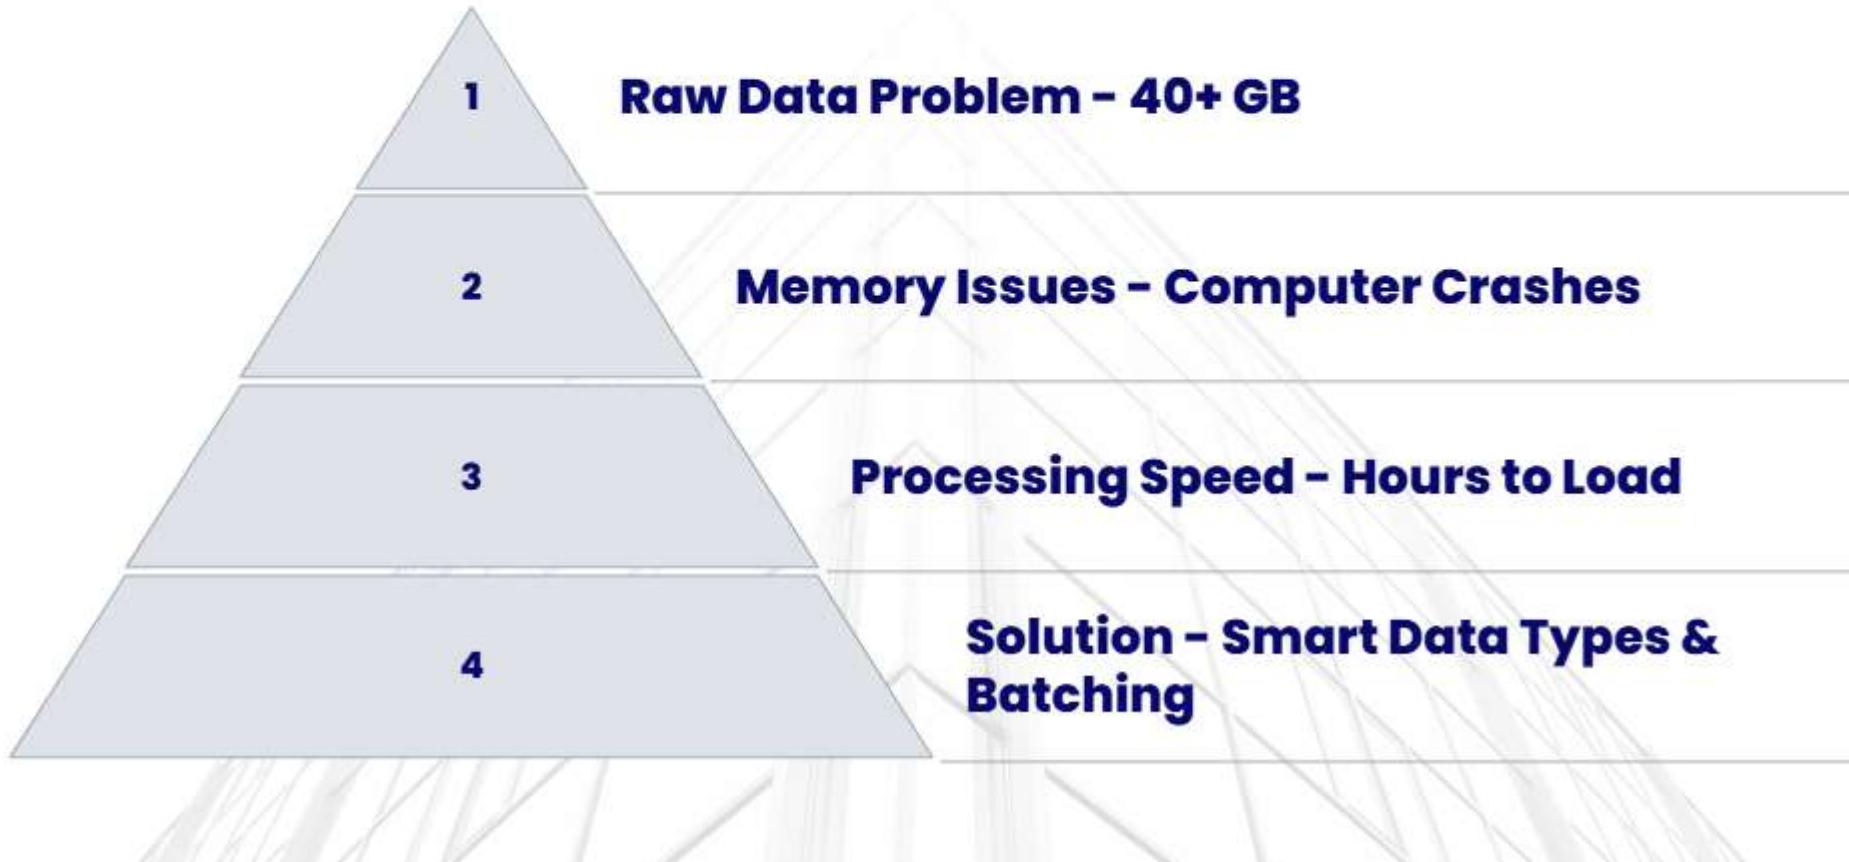


##### Why this matters

Ensures the entire dataset fits in RAM, prevents crashes, and speeds up subsequent feature engineering steps.

# Why Large Retail Data Needs Memory Optimization

## Handling Millions of Transactions Efficiently



# 4. Dataset Linkage Strategy

Merging Disconnected Sources into a Single Timeline

DATA INTEGRATION



## \_JOIN Keys (Logic)

**Sales**      ⚡ Calendar using `day_id`

**Sales**      ⚡ Prices using `item, store, week`

## 💡 Why Linkage is Critical

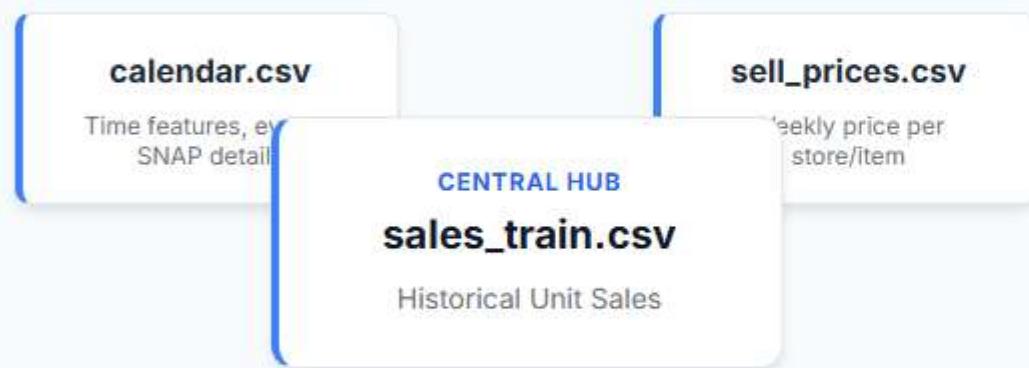
Machine learning models require a **single consistent view**. Linking ensures every "unit sold" is matched with its exact selling price and whether that day was a holiday.

# Dataset Linkage Strategy

DATA PREPROCESSING

Connecting sales, price, and calendar data for forecasting

## Relationship Diagram



### PRIMARY JOIN KEYS

`id` `item_id` `store_id` `state_id` `wm_yr_wk`

## ID Structure Analysis

Every row in the dataset is uniquely identified by a composite string containing hierarchical information.

EXAMPLE ID: FOODS\_1\_018\_CA\_1\_EVALUATION

FOODS	1	018	CA	1	evaluation
Category	Dept	Item	State	Store	Data Split

### Why this structure matters:

- Allows aggregation at multiple levels (State, Store, Category, Department).
- `store_id + item_id` forms the crucial link to pricing data.

# 5. Wide Format vs Long Format Sales Data

DATA TRANSFORMATION

Reshaping data structure for effective machine learning

## ↔ Wide Format (Raw Input)

ITEM_ID	D_1	D_2	...	D_1913
HOBBIES_1	0	0	...	1
HOBBIES_2	0	1	...	0
... 30,000 items ...				

VS

## ↔ Long Format (Target)

ITEM_ID	DAY	SALES
HOBBIES_1	d_1	0
HOBBIES_1	d_2	0
HOBBIES_1	d_3	0
... millions of rows ...		

- ✖ Hard to Link: Cannot easily join "Day 1" with a calendar date column.
- ✖ Bad for ML: Models expect a fixed number of features, but days keep increasing.
- Structure: 1 row per item, many columns.

- ✓ Easy Linkage: "Day" becomes a key to join with Calendar/Prices.
- ✓ ML Ready: Standard "Samples x Features" matrix format.
- ⌚ Structure: 1 row per transaction (Item-Day).

## DECISION:

Convert to Long Format



### Enables Joins

Merges seamlessly with Calendar & Prices



### Time Features

Allows calculating lags (t-1, t-7)



### Model Input

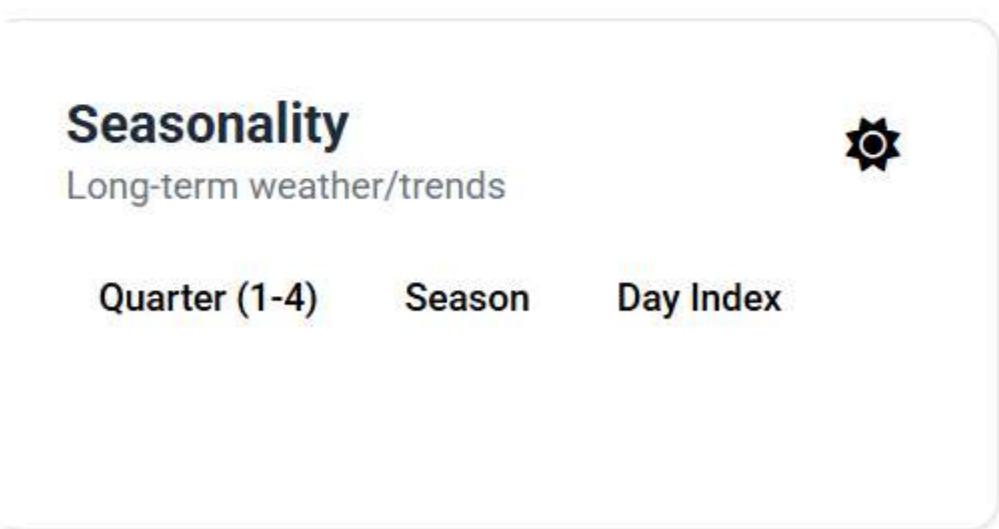
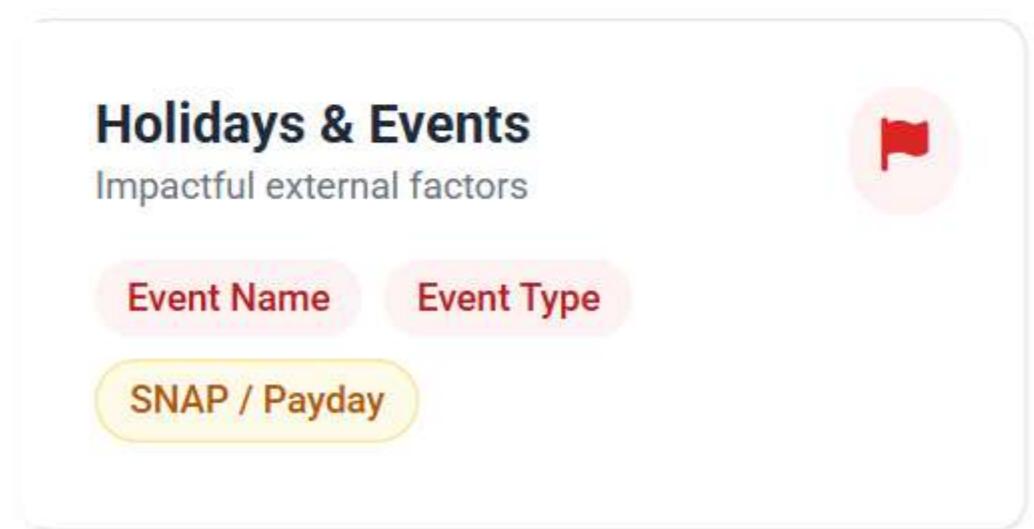
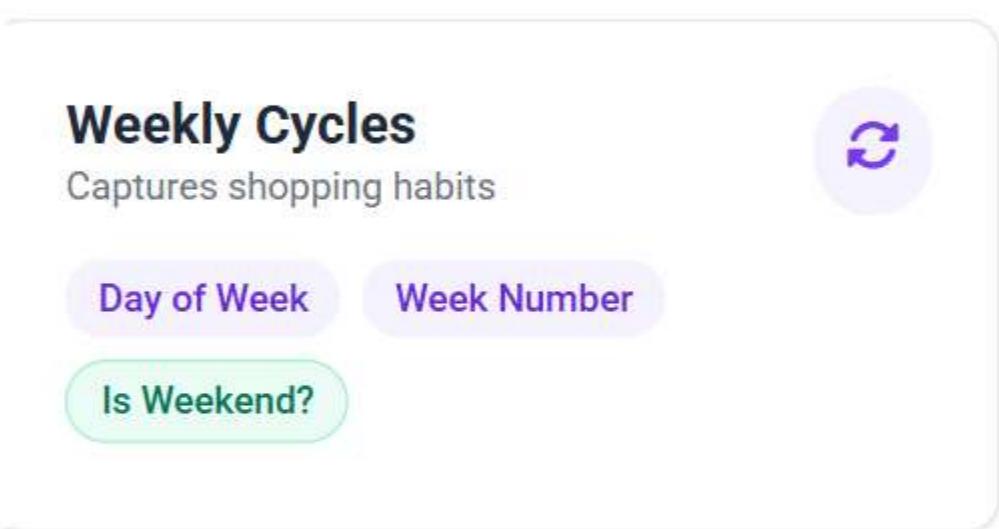
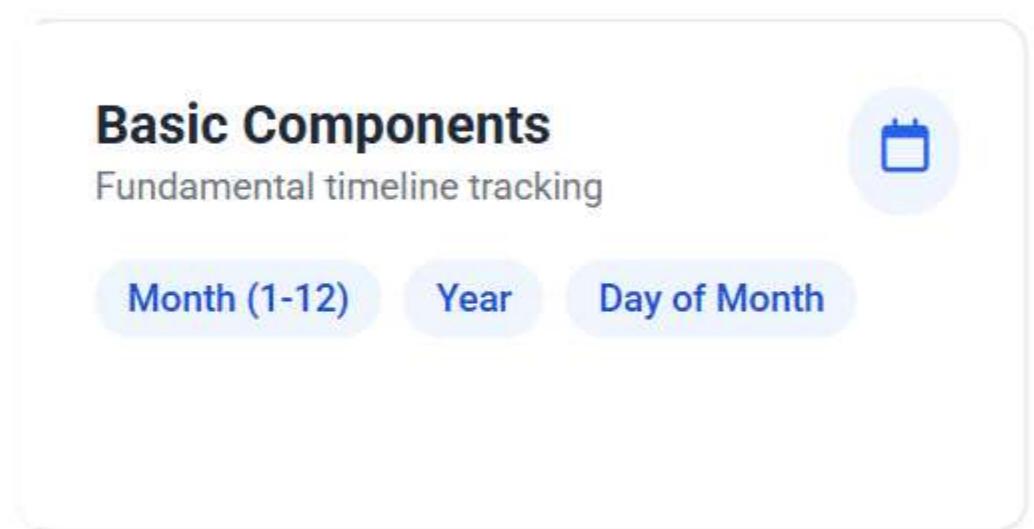
Standard input shape for Forecasting Models

# 6. Time-Based Feature Engineering

## Extracting Demand Signals from Dates

FEATURE ENGINEERING

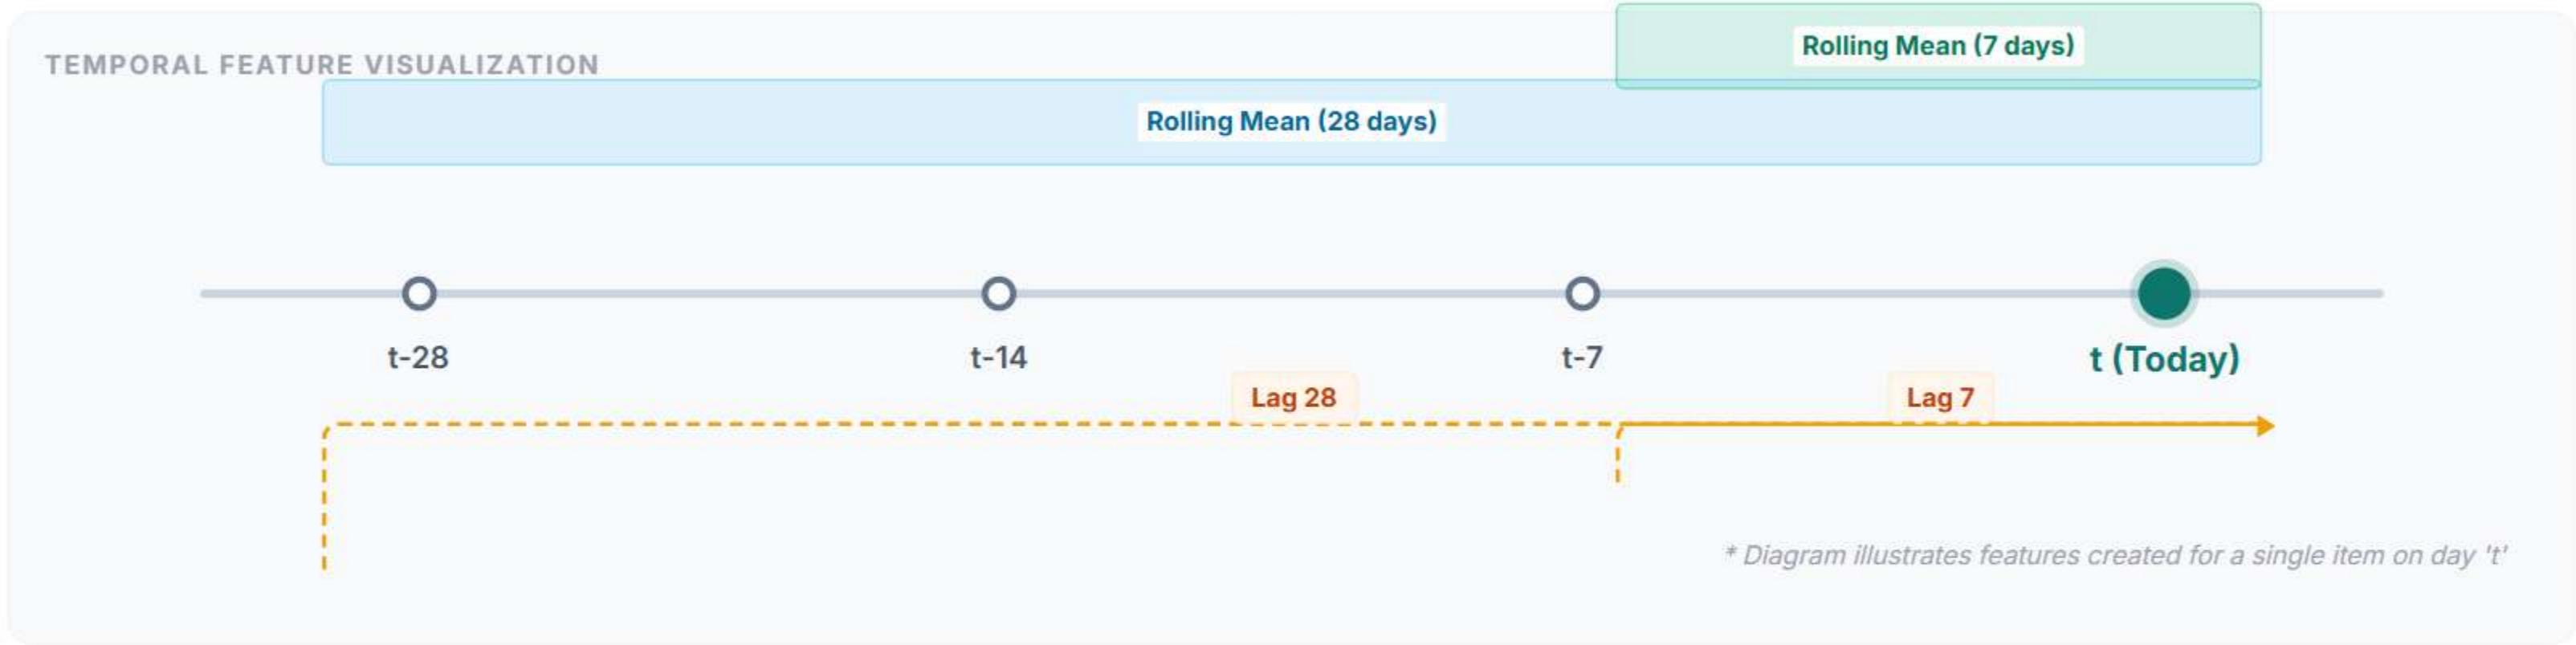
↓ Decomposition of Raw 'Date' Column



- Why it matters**
- ✓ Retail sales follow strict **weekly cycles** (e.g., Saturday peak).
  - ✓ **Holidays** cause massive demand spikes that models miss without explicit flags.
  - ✓ Numerical features allow the model to learn **recurring patterns**.

# 7. Lag and Rolling Window Features

Teaching the model to look back in time



## ⌚ Lag Features

**Definition:** The exact sales value from specific days in the past (e.g., exactly 7 days ago).

**Why it helps:** Captures *seasonality*. If sales are high every Saturday, the "Lag 7" feature tells the model to expect high sales this Saturday.



## Rolling Window Features

**Definition:** Statistical summaries (Mean, Max, Std Dev) over a past window (e.g., average of last 28 days).

**Why it helps:** Captures *recent trends* and stabilizes predictions. If the average is rising, the model predicts an increase.

# Lag and Rolling Window Features

## Learning from Past Sales History

### Step 1

Lag Features - Sales from 1, 7, 28 days ago

### Step 2

Rolling Average - Average sales over past 7, 30 days

### Step 3

Rolling Statistics - Min, max, standard deviation of past sales

# Feature Engineering Overview

Transforming raw time-series data into predictive signals



## Lag Features

- 
- 28 Days
- 35 Days
- 42 Days



## Rolling Means

- 
- 7 Days (Short)
- 28 Days (Medium)
- 60 Days (Long)



## Calendar

- 
- Day of Week
- Month / Year
- Weekend Flag
- SNAP / Holidays



## Price Momentum

- 
- Price Change %
- Max/Min Ratio
- Std. Deviation

*"Captures recent sales volume trends directly"*

*"Smooths volatility to reveal underlying patterns"*

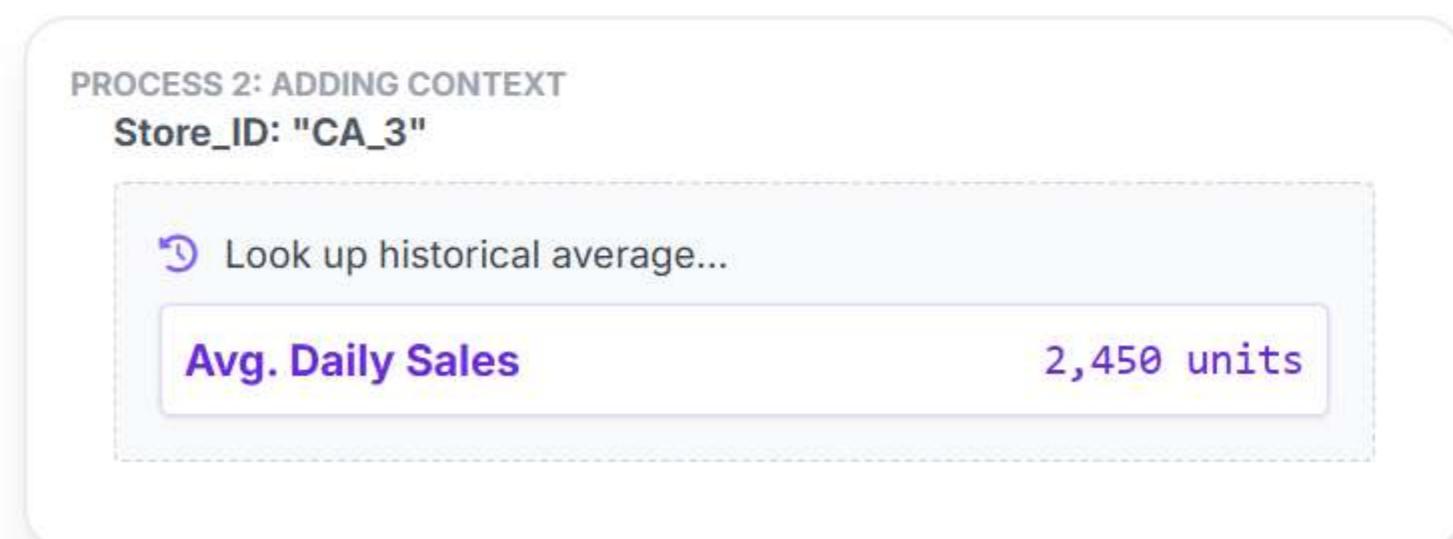
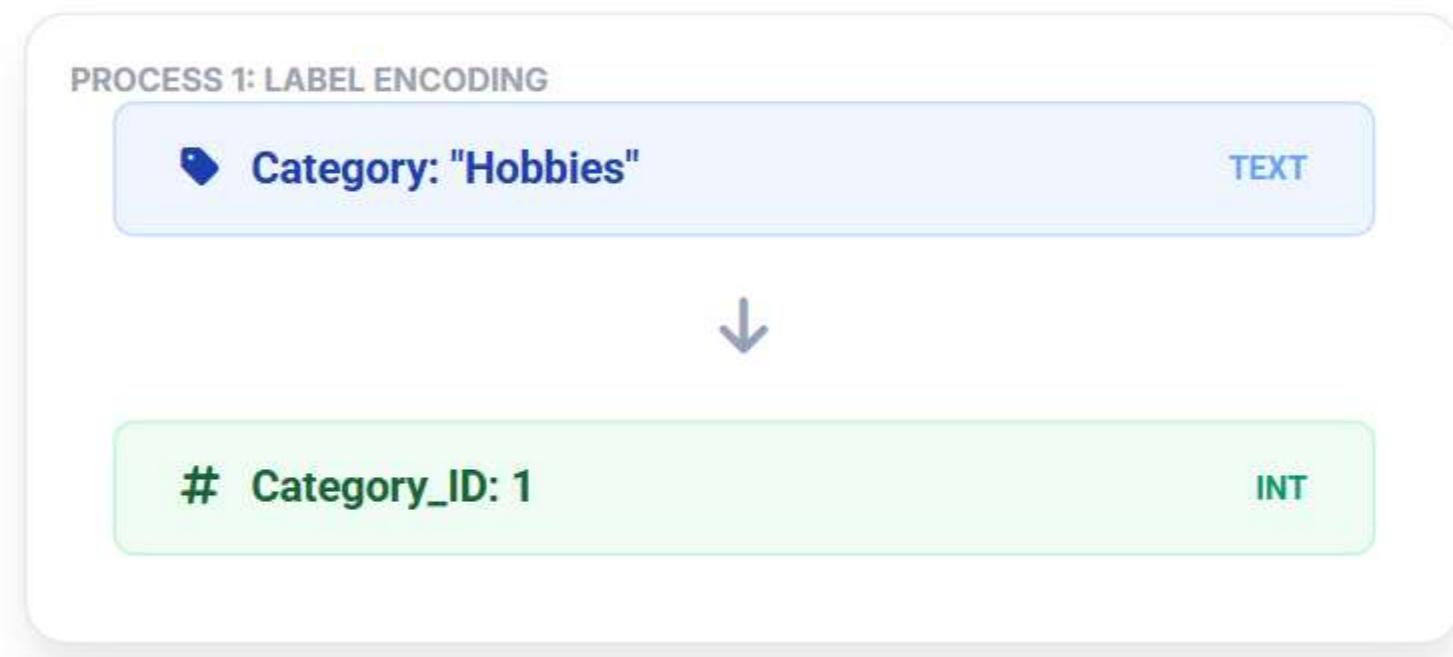
*"Encodes critical seasonality and event markers"*

*"Quantifies price elasticity and discount impact"*

# 8. Handling Categorical Variables

Translating Text Labels into Machine-Readable Signals

DATA PREPARATION



## Why convert to numbers?

Machine Learning models operate on mathematical matrices. They cannot multiply or subtract text strings like "California" or "Foods". We must map every unique text label to a unique integer.



## Grouping adds context

Converting "Store A" to "1" is necessary, but providing the **"Average Sales of Store A"** gives the model a stronger signal about that store's typical performance magnitude.



### CRITICAL PREPROCESSING RULE

When calculating group averages (e.g., average sales per store), use **only past data**. Using future data causes "Data Leakage" and invalidates the model.

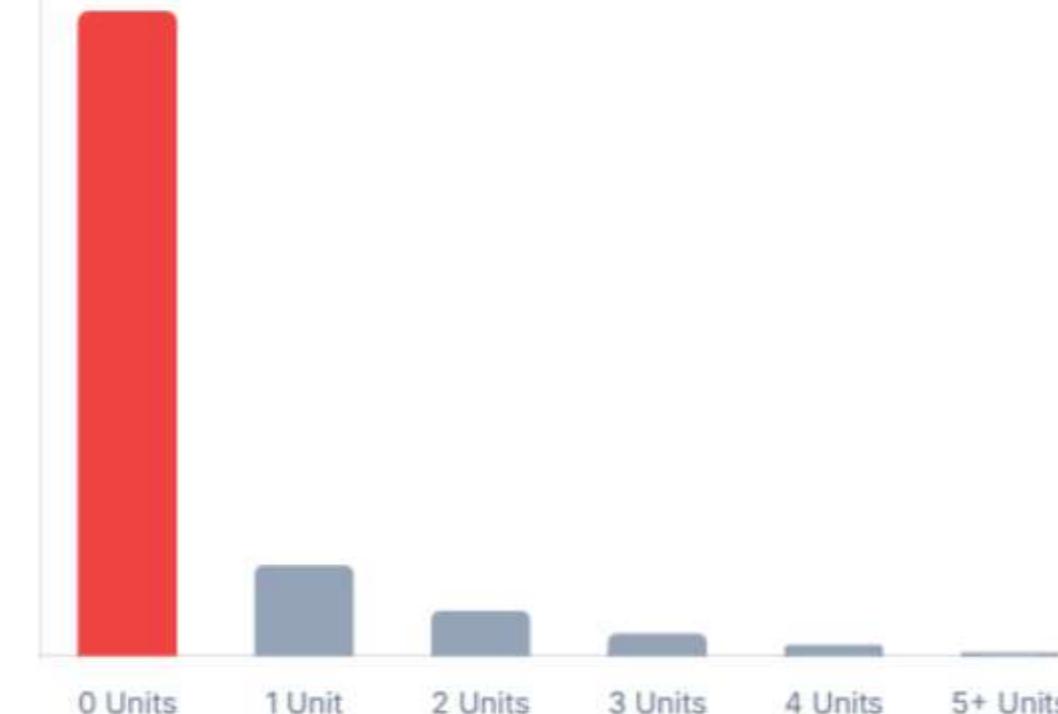
# 9. Handling Zero-Inflated Demand

Solving the "Intermittent Sales" Problem

DATA ENGINEERING

## The Distribution Problem

Frequency of Daily Units Sold



- 💡 **Insight:** Most days have 0 sales. A standard model treats 0 as just a "low number," failing to capture that 0 is a distinct state (No Stock / No Demand).

## ENGINEERED SIGNALS



### Zero Flag (Binary)

Explicitly tells the model: "Was there a sale today?" (0 or 1).



### Recency Gap

"Days since the last non-zero sale." Helps predict when the next sale is due.



### Non-Zero Rate

Rolling probability of a sale occurring (e.g., "Sold on 20% of days last month").



## Why it matters

Without these specific hints, regression models tend to "hedge" their bets and predict impossible values like **0.1 units**. These features help the model separate "Zero Days" from "Sales Days."

# 10. Preventing Data Leakage

Time-Based Splitting Strategy

MODEL VALIDATION



## Split by Time, Not Randomly

**Never shuffle.** Unlike image classification, sales data depends on sequence. We must train on the past to predict the future, preserving the order of events.



## Block Future Information

**No peeking.** Features for a specific day (e.g., "7-day average") must strictly use data *prior* to that day. Using next week's sales to predict today is "leakage."



## Realistic Validation

**Mimic production.** The gap between Training and Validation ensures our model is tested exactly how it will be used in the real world: predicting the unknown.

# 11. Final Processed Dataset

Structure of the Model-Ready Input Matrix

PIPELINE OUTPUT

## ■ SINGLE ROW STRUCTURE (ITEM-STORE-DAY)

IDS & KEYS		TIME FEATURES			PRICE	EVENTS		HISTORY		TARGET
item_id	store_id	day_int	wday	month	sell_price	event_name	snap_CA	lag_7	roll_28	sales (y)
1432	3	1854	6	2	2.48	0	1	5.0	3.2	4
1432	3	1855	7	2	2.48	1	1	4.0	3.1	?

◻ Target Variable (What we predict)

● Categorical Encoded

Engineered Numeric

## Dataset Specs

### GRANULARITY

⌚ Item × Store × Day

### TOTAL FEATURES

⌚ ~35 Columns

### DATA TYPES

⌚ Int16 / Float32

Memory Optimized

### MISSING VALUES

✓ Handled (0 or Imputed)

**MODEL READY**  
VALIDATED FOR TRAINING

# Offline Pre-Aggregation

Optimizing the data pipeline architecture for performance and scalability.



## Join & Clean

Merge raw CSVs and handle missing values.

## Feature Build

Generate lags, rolling means, and calendar features.

## Compress

Downcast types and optimize memory usage.

## Save Dataset

Export optimized pickle for rapid loading.

## Why Offline?

Critical for large-scale retail datasets

### ⚡ Faster Loading

Eliminates repetitive processing time during model training iterations.

### 💾 Reduced RAM

Optimized datatypes significantly lower runtime memory footprint.

### 🎛 Performance

Enables rapid experimentation and smoother application performance.

# Handling Retail-Specific Challenges

Key complexities addressed during data preprocessing

## Intermittent Demand

### Prevalence of Zero Sales Days

Retail data often contains many days with 0 units sold.

Requires specific handling (e.g., Tweedie loss) rather than standard regression.

## Hierarchical Structure

### Multi-Level Aggregation

Data exists at Item → Department → Category → Store → State levels. Trends at aggregate levels may differ from individual item trends.

## Seasonal Patterns

### Event & Holiday Effects

Significant sales spikes driven by events (Super Bowl, Christmas, SNAP days). Mapped via `calendar.csv` features.

## State & Store Variations

### Local Price & Demand Dynamics

Prices and preferences vary by location (CA, TX, WI). 30,490 distinct product-store combinations must be modeled.

# 12. Why This Pipeline Improves Accuracy

From Raw Data to High-Performance Signals

PERFORMANCE IMPACT



## Noise Reduction

Handling missing values and zero-inflated demand prevents the model from learning "garbage" patterns.



## Context Awareness

Correctly linking prices and calendar events explains *why* sales spikes happen (e.g., promotions).



## Stronger Signals

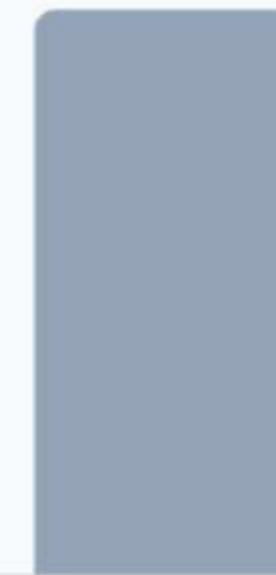
Lags and rolling windows allow the model to "see" trends and seasonality directly.



## Honest Validation

Time-based splitting ensures the accuracy we measure is realistic and reliable for future use.

## FORECAST ERROR (RMSE)



Raw Data Model



With Pipeline



~30% Error Reduction



**Result:** A model-ready dataset that yields higher, more stable accuracy and faster training.

# Final Processed Dataset

## Dataset Contents



### Clean, linked data

Merged from sales, calendar, and pricing sources



### Engineered features

Lag variations (28-42 days) & Rolling windows (7-60 days)



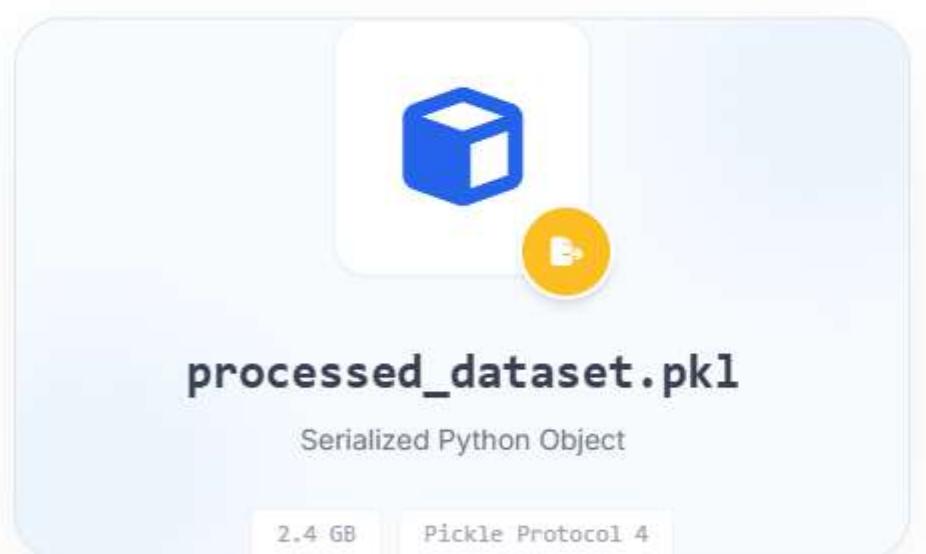
### Encoded categoricals

Label encoding for products, stores, and events



### Memory-optimized dtypes

Downcasted numerics for efficient loading



### STATUS

Ready for ML Model Training

